

The **Core** Standard.

METHODOLOGY · THE FIVE PILLARS AND THE RUBRIC

How we *score*.

The Core Standard is a scenario-based executive assistant assessment: a realistic working day, scored against five pillars on a published, behaviourally anchored rubric, and reviewed by a specialist. This document sets out the framework, the scoring, the reviewer protocol and the safeguards, for anyone who needs to stand over a hiring decision in writing.

Written for HR directors, procurement teams and hiring managers evaluating the assessment.

A realistic working day. *Measured.* Not an interview.

Interviews show how a candidate presents. A CV shows where they have been. Neither shows how they will handle a principal's inbox at 8:45 on a Tuesday. The Core Standard does, by putting every candidate inside the same simulated day and scoring what they actually do.

What it is

An immersive simulation of an Executive Assistant's working day, completed online in a single timed session of 60 to 90 minutes depending on level. The candidate works through a sequence of scenario modules for a fictional principal: an inbox to triage, competing demands, a diary that falls apart, drafting in the principal's voice, a sensitive situation. Every input they write is scored against the rubric and reviewed by a specialist before the report is released. One scenario, not a question bank.

What it is not

- Not a typing test or software quiz. Those measure the parts of the role least likely to go wrong, and they belong to a version of the role that no longer exists.
- Not a psychometric. A psychometric describes traits a candidate reports about themselves. The Core Standard observes behaviour.
- Not multiple choice. Candidates write their decisions and their reasoning. The reasoning is what is scored.
- Not a pass/fail test. There is no pass mark. The report gives pillar scores, an overall band and a narrative, for the employer to weigh alongside interviews and references.

Why it exists

The hires that go wrong rarely go wrong at interview. They interview well, then the diary collapses on a Wednesday and nobody has seen what this person does under that kind of pressure. The Core Standard was built by a specialist EA recruiter to close that gap: to show an employer the work before the offer, and to give EAs a way to prove what interviews never capture.

Two editions. Three levels. *One rubric.*

Employer edition

Commissioned by the hiring organisation for a named candidate. The report goes to the employer, with a reviewer recommendation calibrated to the level being hired.

Candidate edition (TCS Certified)

Sat by an EA on their own initiative. One sitting, scored against every level of the standard, so the result is the level demonstrated. The candidate owns the report, a Recruiter Summary, a digital badge and a certificate. The behavioural anchors are identical across both editions at every level. The scenario changes; the rubric does not. A Level 2 Certified result is directly comparable to a Level 2 employer report.

The three levels

The right level is set by who the assistant supports and how much judgement the seat demands, not by years of experience.

LEVEL	SEAT	WHAT IT TESTS	SESSION
1 · Operational EA	Single-principal support to a department head or senior manager	Inbox triage, diary management, travel and logistics, internal stakeholder communication. Can they hold the basics reliably?	60 min
2 · EA to Senior Leadership	Multi-principal or C-suite support	Sophisticated diary choreography, board logistics, discreet personal matters, cross-functional facilitation. How they read a situation and handle stakeholders under pressure.	75 min
3 · Executive Business Partner	Second-brain support to a C-suite executive or the board	Strategic project oversight, board pack preparation, stakeholder arbitration, delegated decisions within a defined scope.	90 min

Chief of Staff is assessed as a separate stream with its own scenario, rubric and reviewer, because it tests capabilities that do not sit on the EA ladder: proxy decision-making, influence without authority, board-level writing and strategic synthesis.

Five pillars, *one working day*.

Every candidate is scored against the same five pillars at every level. They are dispositions, not skills: the pattern of judgement calls that separates an EA who holds a seat for a decade from one who leaves within the year.

01 **UNDER PRESSURE** **Prioritisation**

Does the candidate recognise what matters most when everything is urgent? Can they say no to the right request at the right time? We score the reasoning behind the ranking, not just the ranking itself.

02 **SIGNAL AND NUANCE** **Judgement**

Do they spot the risk buried in the meeting request, the conflict in the diary, the tone shift in the email? The unflashy work of noticing. Second read, not first glance.

03 **tone and audience** **Communication**

Can they write in the voice of the principal? Concise, warm and appropriate to the audience, whether that is a board member, a supplier or a colleague. We read what they draft, not what they say they would draft.

04 **PLAN TO DELIVERY** **Execution**

Do they close the loop? Do the actions get done to standard, on time, with evidence? Do they manage upward and across, or just react?

05 **TRUST AND PROXY** **Partnership**

Do they work as a genuine partner, not just a pair of hands? Anticipating what the principal needs before it is asked, holding relationships and stakeholders on their behalf, and earning the trust to act in their name.

The dimensions *interviews never reach*.

Higher levels add dimensions suited to the seat. They are scored alongside the five pillars, and they are the ones that most often decide whether a senior hire holds.

Mindset

How the candidate frames the seat, handles friction and thinks about accountability. Not personality: disposition to the work itself. A short layer of five questions runs alongside the modules and is scored separately as a qualitative overlay rather than pass/fail, so the reviewer can see whether the candidate understands the role they are being hired into.

AI orientation

How comfortable and discerning they are using AI tools inside the seat. Where they lean on them, where they override them, and how they judge the output. Assessed at Level 2 and Level 3 through a dedicated module in which the candidate is given an AI-drafted output with a flaw buried in it. An increasingly non-negotiable EA competency, and one that tests judgement rather than tool trivia.

Data interpretation

At Level 3, a further module presents a data set with a decision buried in it: a KPI dashboard or a messy spreadsheet, with the brief to surface the three things the principal needs to know and make a recommendation.

Comparability

Because the rubric is fixed and the scenario is the same for every candidate at a level, scores are comparable across candidates in one search, across searches over time, and between the employer and candidate editions. This is what makes the report a benchmark rather than an opinion.

Five bands, *behaviourally anchored*.

Each pillar is scored on a 0 to 100 benchmark against a published rubric, and each score maps to one of five bands. Every band is anchored to observable behaviour, not adjectives.

Developing

Below the bar for the level

Proficient

Holds the seat

Strong

Reliable under pressure

Distinctive

Top of the level

Exemplary

Ready for the level above

What an anchor looks like

A "Distinctive" on Prioritisation is not "was really good at prioritising". It is, specifically: re-sequenced the day without prompting when a new input arrived; held the principal's agenda in mind when the obvious path would have solved the smaller problem; named the trade-off in writing rather than just picking one.

How scores combine

A candidate's overall band is not an average. It is the lowest pillar, with a minimum floor. Hiring on an average hides the weakness that will cost you most. A candidate who is Distinctive on four pillars and Developing on Partnership is a different hire from one who is Strong on all five, and the report says so.

Red flags and standouts

A red flag is a specific behaviour that would cause real harm in the seat: a confidential detail shared, a principal's commitment made without authority, a sensitive message sent in the wrong register. A red flag caps the relevant pillar at the lowest band for that module, and is named in the report. Standout moves, the decisions a reviewer would want to show a hiring manager, are logged during scoring and carried through to the reviewer narrative.

Every report is reviewed *by an expert.*

Nothing is machine-only. The rubric is applied identically to every candidate, and a specialist reads every submission before a report is released.

The first pass and the second

Every response is scored first against the rubric with AI assistance, applied in the same way to every candidate at the level. A specialist reviewer then reads the full submission, adjusts where the rubric has been applied too literally or too generously, and signs off. The reviewer's sign-off appears on the report.

Who reviews

Every report is reviewed and signed off by Lily Shippen, founder of The Core Standard and a specialist Executive Assistant recruiter of thirteen years, across every stream. The rubric is co-authored with Emily Partridge, a former C-suite Executive Assistant and EA Team Manager.

Calibration

- **Quarterly recalibration.** The reviewer pool recalibrates every quarter against a shared library of anchor responses, so that a "Strong" in March means the same as a "Strong" in September.
- **Inter-rater reliability.** A sample of reports is double-reviewed. Where two reviewers diverge, the divergence is examined and drift is corrected at the next calibration.
- **Pseudonymised scoring.** Reviewers never see the candidate's name, company or CV during scoring. The file is pseudonymised before it reaches the reviewer queue. The candidate is identified by level and scenario only.
- **Sign-off on the page.** Every report carries a reviewer sign-off block, so the employer knows who stands behind the recommendation.

Review and appeal

A candidate can request a review within 14 days of receiving their report. A second reviewer re-scores independently. If the second review differs by more than one band on any pillar, the report is rewritten and reissued at no cost.

One working day, *built from modules.*

The assessment is built from scenario modules, not questions. Each stress-tests a capability that decades of hiring, placing and reviewing EAs tell us is decisive at that level. Nothing is there to pad the time.

MODULE	WHAT THE CANDIDATE FACES	LEVELS
Inbox Triage	A morning inbox of competing demands. What to action, hold, escalate or decline, and in what order.	All
Simultaneous Demands	Three things land at once, all urgent. What gets done first, what waits, and can they explain why.	All
Access and Discretion	A request that tests confidentiality, boundaries, and how they protect the principal and sensitive information.	All
Rebuilding the Day	The diary falls apart mid-morning. Re-sequence under pressure without dropping what matters.	All
Representing the Principal	Draft in the principal's voice across audiences: a board member, a supplier, a colleague.	All
A Sensitive Situation	A delicate interpersonal moment, handled with judgement, tact and an instinct for the principal's reputation.	All
Judgement with AI	An AI-drafted output with a flaw buried in it. Do they catch it, fix it, or ship it?	2 and 3
Reading the Numbers	A data set with a decision buried in it. Interpret, spot what matters, make a recommendation.	3

The six core modules run at every level and are recalibrated at each altitude. A Level 1 Inbox Triage is not a Level 3 Inbox Triage, even though both carry the same name. Level 1 runs the core six across a 60-minute day; Level 2 adds the AI module across 75 minutes; Level 3 adds data interpretation across 90 minutes. The stack has been rebuilt twice since 2024 as the profession has changed.

Module names are shown for illustration. The exact briefs, and the curveballs inside them, are not published and stay unseen until the candidate sits the assessment.

What we *actively suppress*.

A benchmark is only worth having if it is fair. These are the sources of bias we design against, and how.

SOURCE OF BIAS	SAFEGUARD
Name, age, education	Stripped from the submission before the reviewer sees it. The reviewer scores what the candidate did, not where they were educated or how their name reads.
Writing style	The rubric scores decisions, not prose quality. A well-reasoned bullet list at Distinctive band outscores a beautifully written paragraph that misses the call.
Time of day	Reviewers score in batches, not in sequence, so a reviewer's third report of the morning is not compared to their eighth of the afternoon.
Accommodations	Candidates can self-certify an additional 25% on the session clock for neurodivergence, disability or a medical reason, at the pre-checkout stage. No diagnosis or evidence is required. Reviewers do not see which reports had an accommodation applied, and the final band is not adjusted down.

Candidate experience

The assessment is text-based and asynchronous, completed in a quiet environment of the candidate's choosing. There is no timer per question, no live interviewer and no verbal pressure. Candidates are told up front what the simulation involves and how their responses will be used. The exercise is designed to reflect the real role and the value it carries, so that a strong EA is engaged by it rather than put off. Many neurodivergent candidates perform better on a scenario-based assessment than in a live interview.

Scenario integrity

Every commissioned assessment uses a composite scenario held only by candidates and reviewers. Candidates receive a single-use link valid for 14 days and sit the session against a hard-stop clock; the session cannot be paused or re-entered.

What the employer *receives*.

A written competency report within 48 hours of the candidate completing the session, typically 12 to 16 pages, built to sit alongside interview notes and references.

SECTION	CONTENTS
Cover	Candidate, level assessed, commissioning employer, date, and the overall rating on a 0 to 100 benchmark: a calibrated position against EAs assessed at the same level, not a percentage.
Pillar breakdown	Five scores, one per pillar, on the same benchmark, shown as a bar chart because the shape is the insight. Four tight scores and one outlier tells you where to probe; five flat scores is often a safer hire than the headline suggests.
Scenario modules	What each module tested, what the candidate produced, and the reviewer's one-line read on each.
Reviewer note	The narrative: how the candidate made decisions, where they were strongest, where to probe in a final conversation.
Red flags and standouts	Specific behaviours, quoted from the simulation, that would cause harm in the seat or that mark the candidate out.
Recommendation	A hire recommendation calibrated to the level, with the reviewer's sign-off.

Who sees what

For employer assessments, the report goes to the commissioning employer. The candidate sees a short confirmation that their submission has been received. For the candidate edition, the candidate receives their own full report, a Recruiter Summary designed to share, a verified digital badge and a certificate, and decides who sees them.

Using the report

The report is a decision tool, not a screening tool. Most employers commission it for their top three to five shortlisted candidates after first-round conversations, and use it alongside their interview notes and references. The cover page is designed to be shared into a decision thread and understood by the whole panel.

How candidate data *is handled*.

Candidates trust us with their work under pressure. Employers trust us with a hiring decision. Both need to know where the data goes.

STORAGE	UK and EU data centres. All data encrypted in transit and at rest.
ACCESS	Every assessment runs under a single-use link tied to one candidate. Submissions are shared only with the commissioning client.
AI TRAINING	Candidate responses are never used to train models. Scoring uses Anthropic's Claude API under enterprise data handling, which excludes prompts and responses from training by default.
RETENTION	Submissions and reports are retained for 12 months from completion, after which personally identifiable data is deleted and only anonymised benchmark statistics remain. Candidates can request earlier deletion at any time.
LAWFUL BASIS	Legitimate interest for employer-commissioned assessments; consent for candidate-commissioned assessments. Full detail in the privacy policy at thecorestandard.co.uk/privacy.html .
TRANSPARENCY	Candidates are told up front what the simulation involves and how their responses will be used, before they begin.
ENTITY	The Core Standard is a trading name of eavolve Ltd, registered in England and Wales, company number 15678513.

The **Core** Standard.

COMMISSIONING

See a candidate do the job *before you offer it.*

Commissioning is self-serve. Choose the level, pay online, and we invite your candidate the same day. Their report is in your inbox within 48 hours of completion. Current pricing for every level, team benchmarking and agency terms is at thecorestandard.co.uk/pricing.html.

EMPLOYERS thecorestandard.co.uk/ea-hire.html

CANDIDATES thecorestandard.co.uk/certified.html

QUESTIONS hello@thecorestandard.co.uk

The benchmark. Not the average.